

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.formula.api as smf
import statsmodels.api as sm
from statsmodels.graphics.regressionplots import influence_plot
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: df = pd.read_csv("50_Startups.csv")
```

```
In [3]: df.head()
```

Out[3]:

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94

```
In [4]: df.tail()
```

Out[4]:

	R&D Spend	Administration	Marketing Spend	State	Profit
45	1000.23	124153.04	1903.93	New York	64926.08
46	1315.46	115816.21	297114.46	Florida	49490.75
47	0.00	135426.92	0.00	California	42559.73
48	542.05	51743.15	0.00	New York	35673.41
49	0.00	116983.80	45173.06	California	14681.40

```
In [5]: df.isna().sum()
```

Out[5]:

R&D Spend	0
Administration	0
Marketing Spend	0
State	0
Profit	0
dtype:	int64

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   R&D Spend             50 non-null    float64
1   Administration        50 non-null    float64
2   Marketing Spend       50 non-null    float64
3   State                 50 non-null    object
```

```
4 Profit 50 non-null float64
dtypes: float64(4), object(1)
memory usage: 2.1+ KB
```

```
In [7]: from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
```

```
In [8]: df["State"]=label_encoder.fit_transform(df["State"])
```

```
In [9]: df.head(10)
```

```
Out[9]:
```

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	2	192261.83
1	162597.70	151377.59	443898.53	0	191792.06
2	153441.51	101145.55	407934.54	1	191050.39
3	144372.41	118671.85	383199.62	2	182901.99
4	142107.34	91391.77	366168.42	1	166187.94
5	131876.90	99814.71	362861.36	2	156991.12
6	134615.46	147198.87	127716.82	0	156122.51
7	130298.13	145530.06	323876.68	1	155752.60
8	120542.52	148718.95	311613.29	2	152211.77
9	123334.88	108679.17	304981.62	0	149759.96

```
In [10]: numerical_feature = [feature for feature in df.columns if df[feature].dtypes!="O"]
```

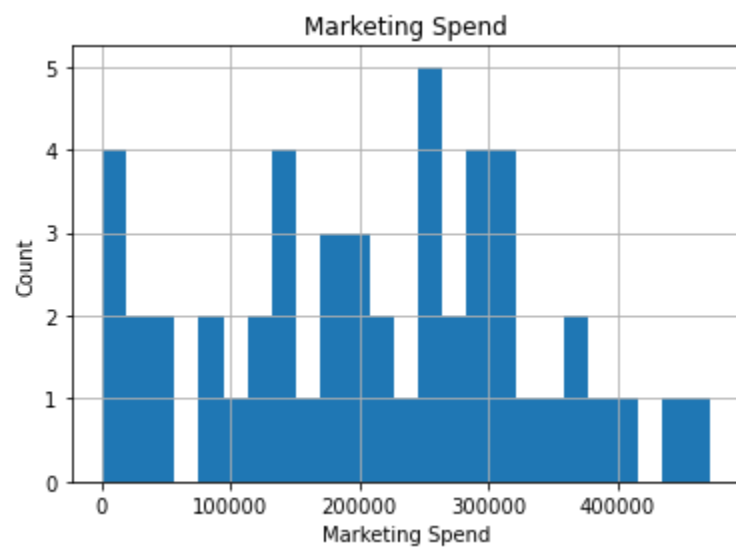

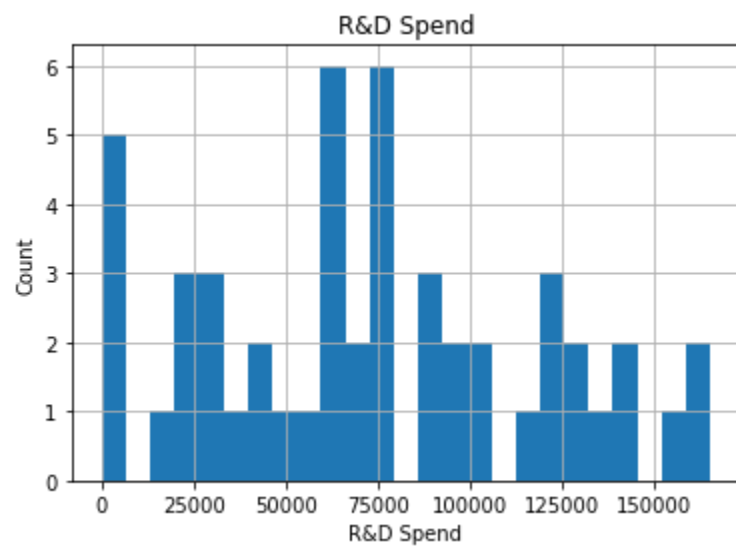
```
In [11]: df[numerical_feature].head()
```


```
Out[11]:
```

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	2	192261.83
1	162597.70	151377.59	443898.53	0	191792.06
2	153441.51	101145.55	407934.54	1	191050.39
3	144372.41	118671.85	383199.62	2	182901.99
4	142107.34	91391.77	366168.42	1	166187.94

Histogram for Raw Data


```
In [13]: for feature in numerical_feature:
data = df.copy()
data[feature].hist(bins=25)
plt.xlabel(feature)
plt.ylabel("Count")
plt.title(feature)
plt.show()
```

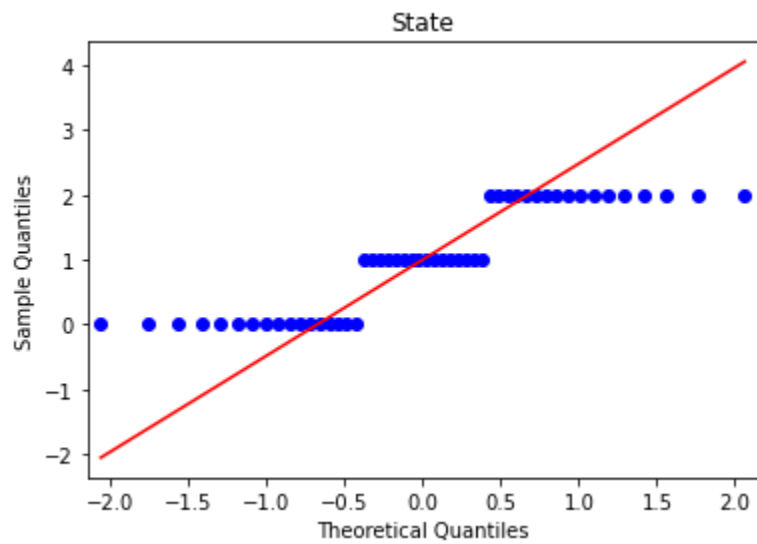






QQ-Plot for Raw Data

```
In [14]: for feature in numerical_feature:
          data = df.copy()
          sm.qqplot(data[feature],line = 'q')
          plt.title(feature)
```












QQ-Plot of Log Transformation

In [14]:



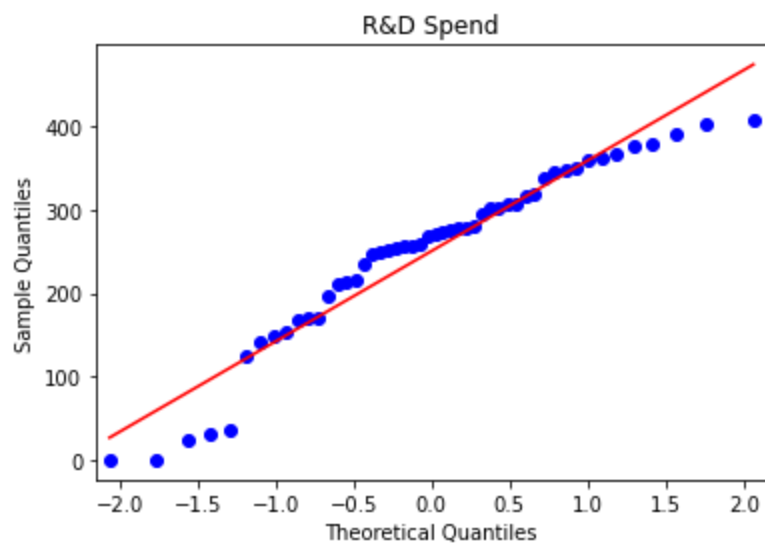
```
for feature in numerical_feature:
    data = df.copy()
    sm.qqplot(np.log(data[feature]), line='q')
    plt.title(feature)
```







QQ-Plot for Square Root Transformation


```
In [15]: for feature in numerical_feature:
        data = df.copy()
        sm.qqplot(np.sqrt(data[feature]), line="r")
        plt.title(feature)
```






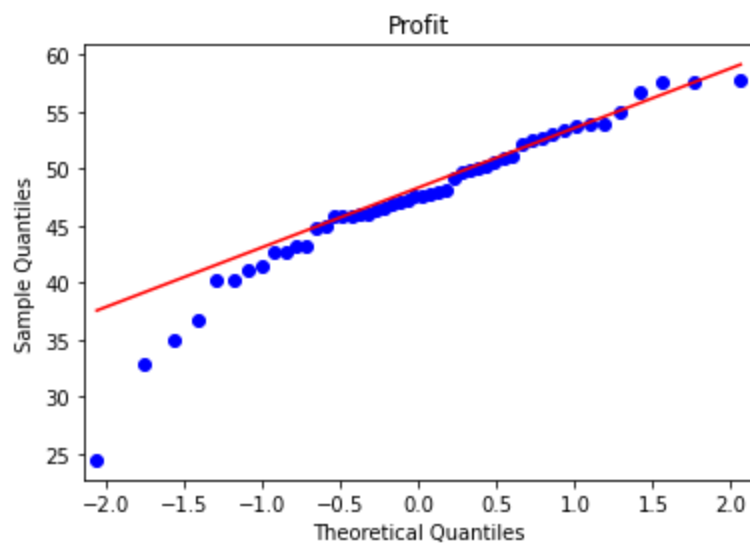


QQ-Plot for Cuberoot Transformation

```
In [16]: for feature in numerical_feature:
data = df.copy()
sm.qqplot(np.cbrt(data[feature]),line="q")
plt.title(feature)
```

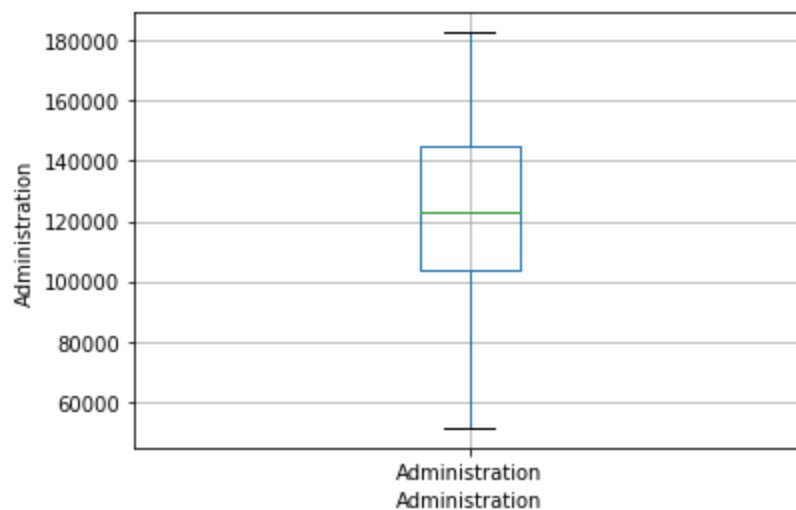



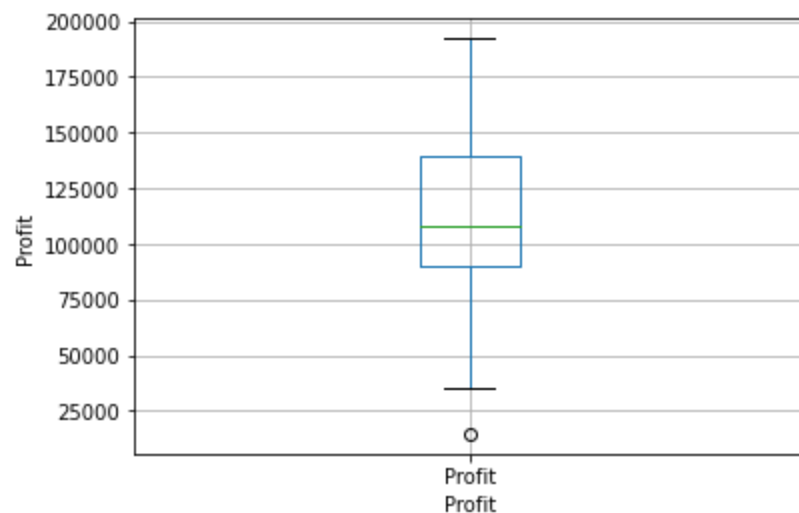

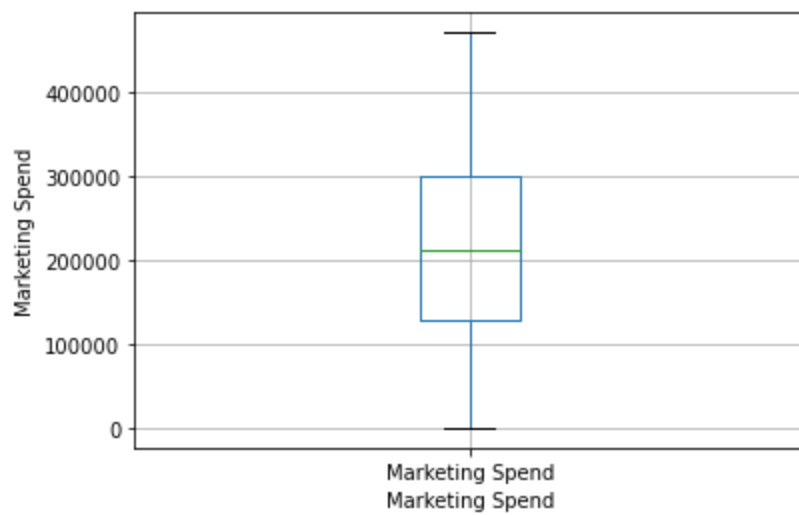




Checking Outliers using Raw Data

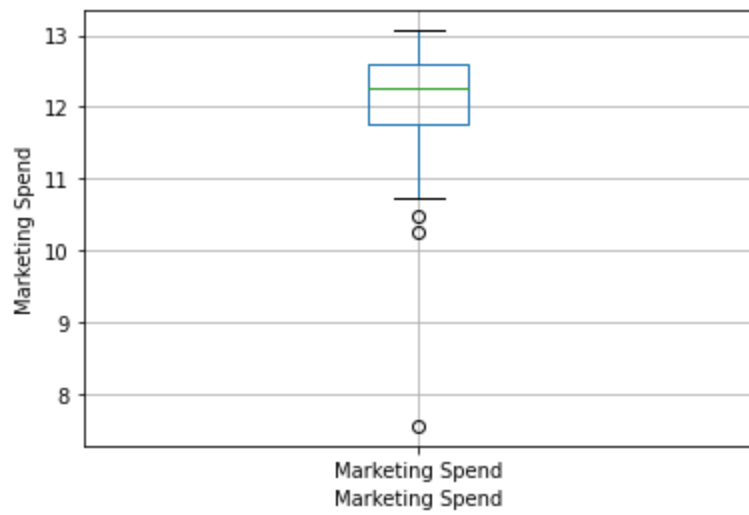


```
In [18]: for feature in numerical_feature:
data = df.copy()
data.boxplot(column=feature)
plt.xlabel(feature)
plt.ylabel(feature)
plt.show()
```







Checking Outliers Using Log Transformation


```
In [17]: for feature in numerical_feature:
          data = df.copy()
          data[feature]=np.log(data[feature])
          data.boxplot(column=feature)
          plt.xlabel(feature)
          plt.ylabel(feature)
          plt.show()
```









Checking Outliers using Squareroot Tarnsformation

```
In [18]: for feature in numerical_feature:
data = df.copy()
data[feature]=np.sqrt(data[feature])
data.boxplot(column=feature)
plt.xlabel(feature)
plt.ylabel(feature)
plt.show()
```












Checking Outliers using Cuberoot Transformation

```
In [19]: for feature in numerical_feature:
data = df.copy()
data[feature]=np.cbrt(data[feature])
data.boxplot(column=feature)
plt.xlabel(feature)
plt.ylabel(feature)
plt.show()
```





Observation:

Hence we see that the raw data having less outliers and as a result they are normally distributed as well but we do transformation in the data we can see the outliers and the little bit skewness in the data. so we choose the raw data because it is normally distributed and having less outliers compared to transformation and also the state column cannot be usefull in this data set so we can drop the state column because it cannot contributing much in the data set for predicting

Checking Co-linearity in the Data

In [20]:

```
df.corr()
```

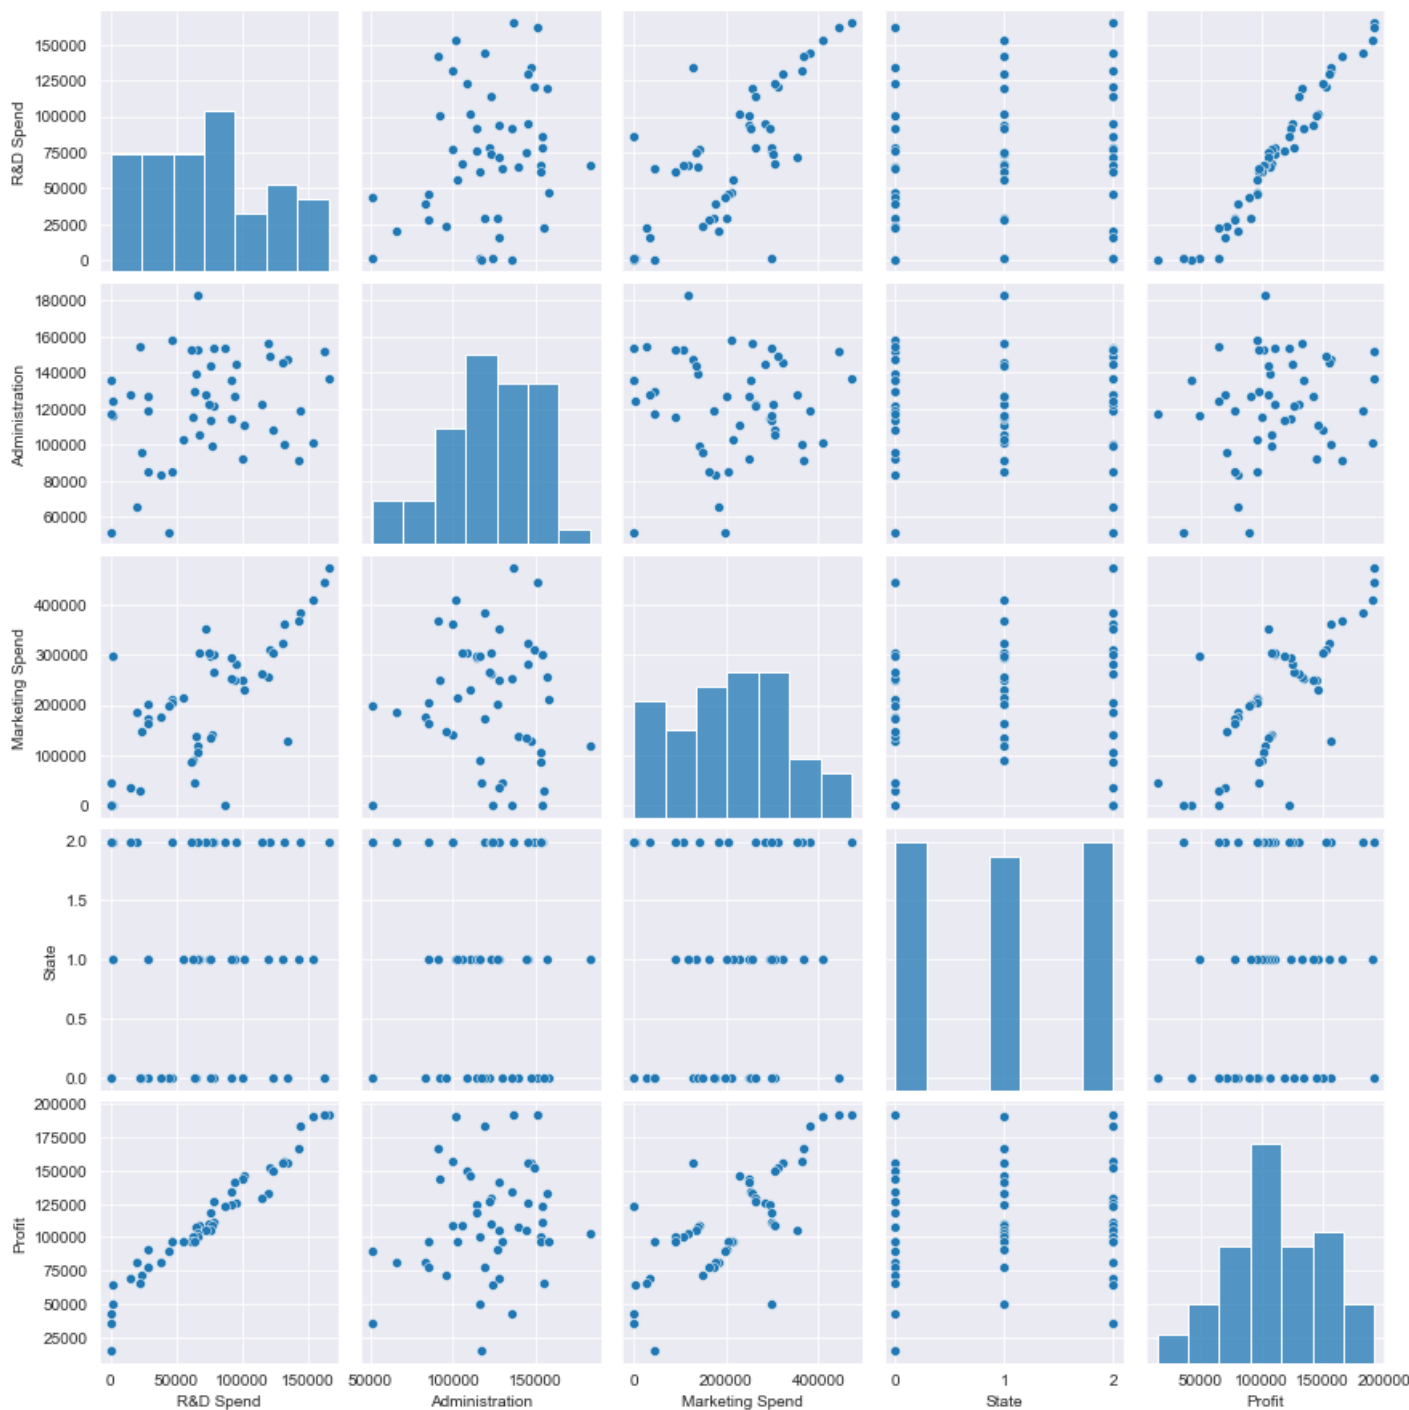
Out[20]:

	R&D Spend	Administration	Marketing Spend	State	Profit
R&D Spend	1.000000	0.241955	0.724248	0.104685	0.972900
Administration	0.241955	1.000000	-0.032154	0.011847	0.200717
Marketing Spend	0.724248	-0.032154	1.000000	0.077670	0.747766
State	0.104685	0.011847	0.077670	1.000000	0.101796
Profit	0.972900	0.200717	0.747766	0.101796	1.000000

In [21]:

```
sns.set_style(style = 'darkgrid')  
sns.pairplot(df)
```

Out[21]: <seaborn.axisgrid.PairGrid at 0x26267e9dca0>



Observation:

As we can see that the state is not contributing as much in the dataset so we can remove state column.R&D and Profit are higly correalted.R&D and Marketing also correlated|

```
In [22]: df.drop("State",axis=1,inplace=True)
```

```
In [23]: df
```

Out[23]:		R&D Spend	Administration	Marketing Spend	Profit
	0	165349.20	136897.80	471784.10	192261.83
	1	162597.70	151377.59	443898.53	191792.06
	2	153441.51	101145.55	407934.54	191050.39
	3	144372.41	118671.85	383199.62	182901.99
	4	142107.34	91391.77	366168.42	166187.94
	5	131876.90	99814.71	362861.36	156991.12
	6	134615.46	147198.87	127716.82	156122.51
	7	130298.13	145530.06	323876.68	155752.60
	8	120542.52	148718.95	311613.29	152211.77
	9	123334.88	108679.17	304981.62	149759.96
	10	101913.08	110594.11	229160.95	146121.95
	11	100671.96	91790.61	249744.55	144259.40
	12	93863.75	127320.38	249839.44	141585.52
	13	91992.39	135495.07	252664.93	134307.35
	14	119943.24	156547.42	256512.92	132602.65
	15	114523.61	122616.84	261776.23	129917.04
	16	78013.11	121597.55	264346.06	126992.93
	17	94657.16	145077.58	282574.31	125370.37
	18	91749.16	114175.79	294919.57	124266.90
	19	86419.70	153514.11	0.00	122776.86
	20	76253.86	113867.30	298664.47	118474.03
	21	78389.47	153773.43	299737.29	111313.02
	22	73994.56	122782.75	303319.26	110352.25
	23	67532.53	105751.03	304768.73	108733.99
	24	77044.01	99281.34	140574.81	108552.04
	25	64664.71	139553.16	137962.62	107404.34
	26	75328.87	144135.98	134050.07	105733.54
	27	72107.60	127864.55	353183.81	105008.31
	28	66051.52	182645.56	118148.20	103282.38
	29	65605.48	153032.06	107138.38	101004.64
	30	61994.48	115641.28	91131.24	99937.59

	R&D Spend	Administration	Marketing Spend	Profit
31	61136.38	152701.92	88218.23	97483.56
32	63408.86	129219.61	46085.25	97427.84
33	55493.95	103057.49	214634.81	96778.92
34	46426.07	157693.92	210797.67	96712.80
35	46014.02	85047.44	205517.64	96479.51
36	28663.76	127056.21	201126.82	90708.19
37	44069.95	51283.14	197029.42	89949.14
38	20229.59	65947.93	185265.10	81229.06
39	38558.51	82982.09	174999.30	81005.76
40	28754.33	118546.05	172795.67	78239.91
41	27892.92	84710.77	164470.71	77798.83
42	23640.93	96189.63	148001.11	71498.49
43	15505.73	127382.30	35534.17	69758.98
44	22177.74	154806.14	28334.72	65200.33
45	1000.23	124153.04	1903.93	64926.08
46	1315.46	115816.21	297114.46	49490.75
47	0.00	135426.92	0.00	42559.73
48	542.05	51743.15	0.00	35673.41
49	0.00	116983.80	45173.06	14681.40

In [24]:

df1 = df.rename({"R&D Spend":"RDS","Marketing Spend":"Marketing"},axis=1)
df1

Out[24]:

	RDS	Administration	Marketing	Profit
0	165349.20	136897.80	471784.10	192261.83
1	162597.70	151377.59	443898.53	191792.06
2	153441.51	101145.55	407934.54	191050.39
3	144372.41	118671.85	383199.62	182901.99
4	142107.34	91391.77	366168.42	166187.94
5	131876.90	99814.71	362861.36	156991.12
6	134615.46	147198.87	127716.82	156122.51
7	130298.13	145530.06	323876.68	155752.60
8	120542.52	148718.95	311613.29	152211.77
9	123334.88	108679.17	304981.62	149759.96
10	101913.08	110594.11	229160.95	146121.95
11	100671.96	91790.61	249744.55	144259.40
12	93863.75	127320.38	249839.44	141585.52
13	91992.39	135495.07	252664.93	134307.35
14	119943.24	156547.42	256512.92	132602.65
		122616.84	261776.23	129917.04

	RDS	Administration	Marketing	Profit
16	78013.11	121597.55	264346.06	126992.93
17	94657.16	145077.58	282574.31	125370.37
18	91749.16	114175.79	294919.57	124266.90
19	86419.70	153514.11	0.00	122776.86
20	76253.86	113867.30	298664.47	118474.03
21	78389.47	153773.43	299737.29	111313.02
22	73994.56	122782.75	303319.26	110352.25
23	67532.53	105751.03	304768.73	108733.99
24	77044.01	99281.34	140574.81	108552.04
25	64664.71	139553.16	137962.62	107404.34
26	75328.87	144135.98	134050.07	105733.54
27	72107.60	127864.55	353183.81	105008.31
28	66051.52	182645.56	118148.20	103282.38
29	65605.48	153032.06	107138.38	101004.64
30	61994.48	115641.28	91131.24	99937.59
31	61136.38	152701.92	88218.23	97483.56
32	63408.86	129219.61	46085.25	97427.84
33	55493.95	103057.49	214634.81	96778.92
34	46426.07	157693.92	210797.67	96712.80
35	46014.02	85047.44	205517.64	96479.51
36	28663.76	127056.21	201126.82	90708.19
37	44069.95	51283.14	197029.42	89949.14
38	20229.59	65947.93	185265.10	81229.06
39	38558.51	82982.09	174999.30	81005.76
40	28754.33	118546.05	172795.67	78239.91
41	27892.92	84710.77	164470.71	77798.83
42	23640.93	96189.63	148001.11	71498.49
43	15505.73	127382.30	35534.17	69758.98
44	22177.74	154806.14	28334.72	65200.33
45	1000.23	124153.04	1903.93	64926.08
46	1315.46	115816.21	297114.46	49490.75
47	0.00	135426.92	0.00	42559.73
48	542.05	51743.15	0.00	35673.41
49	0.00	116983.80	45173.06	14681.40

Creating the first model

```
In [25]: model1 = smf.ols("Profit~RDS+Administration+Marketing",data=df1).fit()
```

In [26]: `model1.summary()`

Out[26]:

OLS Regression Results							
Dep. Variable:		Profit		R-squared:		0.951	
Model:		OLS		Adj. R-squared:		0.948	
Method:		Least Squares		F-statistic:		296.0	
Date:		Sun, 10 Apr 2022		Prob (F-statistic):		4.53e-30	
Time:		19:47:24		Log-Likelihood:		-525.39	
No. Observations:		50		AIC:		1059.	
Df Residuals:		46		BIC:		1066.	
Df Model:		3					
Covariance Type:		nonrobust					
		coef	std err	t	P> t	[0.025	0.975]
Intercept		5.012e+04	6572.353	7.626	0.000	3.69e+04	6.34e+04
RDS		0.8057	0.045	17.846	0.000	0.715	0.897
Administration		-0.0268	0.051	-0.526	0.602	-0.130	0.076
Marketing		0.0272	0.016	1.655	0.105	-0.006	0.060
Omnibus:		14.838	Durbin-Watson:		1.282		
Prob(Omnibus):		0.001	Jarque-Bera (JB):		21.442		
Skew:		-0.949	Prob(JB):		2.21e-05		
Kurtosis:		5.586	Cond. No.		1.40e+06		

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.4e+06. This might indicate that there are strong multicollinearity or other numerical problems.

In [27]: `model1.params`

Out[27]:

Intercept	50122.192990
RDS	0.805715
Administration	-0.026816
Marketing	0.027228

dtype: float64

Simple Linear Regression

In [28]: `model2 = smf.ols("Profit~Administration",data=df1).fit()`

In [29]: `model2.summary()`

Out[29]:

OLS Regression Results			
Dep. Variable:	Profit	R-squared:	0.040

Model:	OLS	Adj. R-squared:	0.020
Method:	Least Squares	F-statistic:	2.015
Date:	Sun, 10 Apr 2022	Prob (F-statistic):	0.162
Time:	19:47:27	Log-Likelihood:	-599.63
No. Observations:	50	AIC:	1203.
Df Residuals:	48	BIC:	1207.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	7.697e+04	2.53e+04	3.040	0.004	2.61e+04	1.28e+05
Administration	0.2887	0.203	1.419	0.162	-0.120	0.698

Omnibus:	0.126	Durbin-Watson:	0.099
Prob(Omnibus):	0.939	Jarque-Bera (JB):	0.110
Skew:	0.093	Prob(JB):	0.947
Kurtosis:	2.866	Cond. No.	5.59e+05

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.59e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [30]: model3 = smf.ols("Profit~Marketing",data=df1).fit()
model3.summary()
```

OLS Regression Results			
Dep. Variable:	Profit	R-squared:	0.559
Model:	OLS	Adj. R-squared:	0.550
Method:	Least Squares	F-statistic:	60.88
Date:	Sun, 10 Apr 2022	Prob (F-statistic):	4.38e-10
Time:	19:47:27	Log-Likelihood:	-580.18
No. Observations:	50	AIC:	1164.
Df Residuals:	48	BIC:	1168.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	6e+04	7684.530	7.808	0.000	4.46e+04	7.55e+04
Marketing	0.2465	0.032	7.803	0.000	0.183	0.310

Omnibus:	4.420	Durbin-Watson:	1.178
Prob(Omnibus):	0.110	Jarque-Bera (JB):	3.882

Skew: -0.336 Prob(JB): 0.144

Kurtosis: 4.188 Cond. No. 4.89e+05

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.89e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [31]: model4 = smf.ols("Profit~RDS", data=df1).fit()  
model4.summary()
```

```
Out[31]:
```

OLS Regression Results						
Dep. Variable:	Profit	R-squared:	0.947			
Model:	OLS	Adj. R-squared:	0.945			
Method:	Least Squares	F-statistic:	849.8			
Date:	Sun, 10 Apr 2022	Prob (F-statistic):	3.50e-32			
Time:	19:47:28	Log-Likelihood:	-527.44			
No. Observations:	50	AIC:	1059.			
Df Residuals:	48	BIC:	1063.			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	4.903e+04	2537.897	19.320	0.000	4.39e+04	5.41e+04
RDS	0.8543	0.029	29.151	0.000	0.795	0.913
Omnibus:	13.727	Durbin-Watson:	1.116			
Prob(Omnibus):	0.001	Jarque-Bera (JB):	18.536			
Skew:	-0.911	Prob(JB):	9.44e-05			
Kurtosis:	5.361	Cond. No.	1.65e+05			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.65e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [32]: model5 = smf.ols("Profit~Administration+Marketing", data=df1).fit()  
model5.summary()
```

```
Out[32]:
```

OLS Regression Results			
Dep. Variable:	Profit	R-squared:	0.610
Model:	OLS	Adj. R-squared:	0.593
Method:	Least Squares	F-statistic:	36.71

Date:	Sun, 10 Apr 2022	Prob (F-statistic):	2.50e-10
Time:	19:47:28	Log-Likelihood:	-577.13
No. Observations:	50	AIC:	1160.
Df Residuals:	47	BIC:	1166.
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.022e+04	1.77e+04	1.143	0.259	-1.54e+04	5.58e+04
Administration	0.3237	0.131	2.468	0.017	0.060	0.588
Marketing	0.2488	0.030	8.281	0.000	0.188	0.309

Omnibus:	6.584	Durbin-Watson:	1.279
Prob(Omnibus):	0.037	Jarque-Bera (JB):	6.524
Skew:	-0.512	Prob(JB):	0.0383
Kurtosis:	4.443	Cond. No.	1.30e+06

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.3e+06. This might indicate that there are strong multicollinearity or other numerical problems.

Transforming Data InTo Standard Scaler for Better Results

```
In [33]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
```

```
In [34]: df2=sc.fit_transform(df1)
df2
```

```
Out[34]: array([[ 2.01641149e+00,  5.60752915e-01,  2.15394309e+00,
  2.01120333e+00],
 [ 1.95586034e+00,  1.08280658e+00,  1.92360040e+00,
  1.99942997e+00],
 [ 1.75436374e+00, -7.28257028e-01,  1.62652767e+00,
  1.98084225e+00],
 [ 1.55478369e+00, -9.63646307e-02,  1.42221024e+00,
  1.77662724e+00],
 [ 1.50493720e+00, -1.07991935e+00,  1.28152771e+00,
  1.35774012e+00],
 [ 1.27980001e+00, -7.76239071e-01,  1.25421046e+00,
  1.12724963e+00],
 [ 1.34006641e+00,  9.32147208e-01, -6.88149930e-01,
  1.10548055e+00],
 [ 1.24505666e+00,  8.71980011e-01,  9.32185978e-01,
  1.09620987e+00],
 [ 1.03036886e+00,  9.86952101e-01,  8.30886909e-01,
  1.00746967e+00],
 [ 8.1921e+00, -4.56640246e-01,  7.76107440e-01,
```

9.46022467e-01],
 [6.20398248e-01, -3.87599089e-01, 1.49807267e-01,
 8.54846746e-01],
 [5.93085418e-01, -1.06553960e+00, 3.19833623e-01,
 8.08167561e-01],
 [4.43259872e-01, 2.15449064e-01, 3.20617441e-01,
 7.41154844e-01],
 [4.02077603e-01, 5.10178953e-01, 3.43956788e-01,
 5.58749518e-01],
 [1.01718075e+00, 1.26919939e+00, 3.75742273e-01,
 5.16026367e-01],
 [8.97913123e-01, 4.58678535e-02, 4.19218702e-01,
 4.48719672e-01],
 [9.44411957e-02, 9.11841968e-03, 4.40446224e-01,
 3.75435696e-01],
 [4.60720127e-01, 8.55666318e-01, 5.91016724e-01,
 3.34771135e-01],
 [3.96724938e-01, -2.58465367e-01, 6.92992062e-01,
 3.07115996e-01],
 [2.79441650e-01, 1.15983657e+00, -1.74312698e+00,
 2.69772649e-01],
 [5.57260867e-02, -2.69587651e-01, 7.23925995e-01,
 1.61935224e-01],
 [1.02723599e-01, 1.16918609e+00, 7.32787791e-01,
 -1.75338400e-02],
 [6.00657792e-03, 5.18495648e-02, 7.62375876e-01,
 -4.16126351e-02],
 [-1.36200724e-01, -5.62211268e-01, 7.74348908e-01,
 -8.21694292e-02],
 [7.31146008e-02, -7.95469167e-01, -5.81939297e-01,
 -8.67294558e-02],
 [-1.99311688e-01, 6.56489139e-01, -6.03516725e-01,
 -1.15493086e-01],
 [3.53702028e-02, 8.21717916e-01, -6.35835495e-01,
 -1.57366637e-01],
 [-3.55189938e-02, 2.35068543e-01, 1.17427116e+00,
 -1.75542334e-01],
 [-1.68792717e-01, 2.21014050e+00, -7.67189437e-01,
 -2.18797551e-01],
 [-1.78608540e-01, 1.14245677e+00, -8.58133663e-01,
 -2.75882217e-01],
 [-2.58074369e-01, -2.05628659e-01, -9.90357166e-01,
 -3.02624599e-01],
 [-2.76958231e-01, 1.13055391e+00, -1.01441945e+00,
 -3.64127442e-01],
 [-2.26948675e-01, 2.83923813e-01, -1.36244978e+00,
 -3.65523895e-01],
 [-4.01128925e-01, -6.59324033e-01, 2.98172434e-02,
 -3.81787113e-01],
 [-6.00682122e-01, 1.31053525e+00, -1.87861793e-03,
 -3.83444211e-01],
 [-6.09749941e-01, -1.30865753e+00, -4.54931587e-02,
 -3.89290919e-01],
 [-9.91570153e-01, 2.05924691e-01, -8.17625734e-02,
 -5.33931605e-01],
 [-6.52532310e-01, -2.52599402e+00, -1.15608256e-01,
 -5.52954899e-01],
 [-1.17717755e+00, -1.99727037e+00, -2.12784866e-01,
 -7.71497339e-01],
 [-7.73820359e-01, -1.38312156e+00, -2.97583276e-01,
 -7.77093678e-01],
 [-9.89577015e-01, -1.00900218e-01, -3.15785883e-01,
 -8.46411346e-01],
 [-1.00853372e+00, -1.32079581e+00, -3.84552407e-01,
 -8.57465682e-01],
 [-1.10210556e+00, -9.06937535e-01, -5.20595959e-01,
 -1.01536466e+00],
 [-1.28113364e+00, 2.17681524e-01, -1.44960468e+00,
 -1.05896021e+00],
 [-3.0539e+00, 1.20641936e+00, -1.50907418e+00,

```

-1.17320899e+00],
[-1.60035036e+00, 1.01253936e-01, -1.72739998e+00,
-1.18008224e+00],
[-1.59341322e+00, -1.99321741e-01, 7.11122474e-01,
-1.56692212e+00],
[-1.62236202e+00, 5.07721876e-01, -1.74312698e+00,
-1.74062718e+00],
[-1.61043334e+00, -2.50940884e+00, -1.74312698e+00,
-1.91321197e+00],
[-1.62236202e+00, -1.57225506e-01, -1.36998473e+00,
-2.43931323e+00]])

```

Again Building Model in Standard Scaler

```

In [35]: data = pd.DataFrame(df2, columns=["RDS", "Administration", "Marketing", "Profit"])
data

```

```

Out[35]:

```

	RDS	Administration	Marketing	Profit
0	2.016411	0.560753	2.153943	2.011203
1	1.955860	1.082807	1.923600	1.999430
2	1.754364	-0.728257	1.626528	1.980842
3	1.554784	-0.096365	1.422210	1.776627
4	1.504937	-1.079919	1.281528	1.357740
5	1.279800	-0.776239	1.254210	1.127250
6	1.340066	0.932147	-0.688150	1.105481
7	1.245057	0.871980	0.932186	1.096210
8	1.030369	0.986952	0.830887	1.007470
9	1.091819	-0.456640	0.776107	0.946022
10	0.620398	-0.387599	0.149807	0.854847
11	0.593085	-1.065540	0.319834	0.808168
12	0.443260	0.215449	0.320617	0.741155
13	0.402078	0.510179	0.343957	0.558750
14	1.017181	1.269199	0.375742	0.516026
15	0.897913	0.045868	0.419219	0.448720
16	0.094441	0.009118	0.440446	0.375436
17	0.460720	0.855666	0.591017	0.334771
18	0.396725	-0.258465	0.692992	0.307116
19	0.279442	1.159837	-1.743127	0.269773
20	0.055726	-0.269588	0.723926	0.161935
21	0.102724	1.169186	0.732788	-0.017534
22	0.006007	0.051850	0.762376	-0.041613
23	-0.136201	-0.562211	0.774349	-0.082169
24	0.073115	-0.795469	-0.581939	-0.086729
25	-0.199312	0.656489	-0.603517	-0.115493
26	0.035370	0.821718	-0.635835	-0.157367

	RDS	Administration	Marketing	Profit
28	-0.168793	2.210141	-0.767189	-0.218798
29	-0.178609	1.142457	-0.858134	-0.275882
30	-0.258074	-0.205629	-0.990357	-0.302625
31	-0.276958	1.130554	-1.014419	-0.364127
32	-0.226949	0.283924	-1.362450	-0.365524
33	-0.401129	-0.659324	0.029817	-0.381787
34	-0.600682	1.310535	-0.001879	-0.383444
35	-0.609750	-1.308658	-0.045493	-0.389291
36	-0.991570	0.205925	-0.081763	-0.533932
37	-0.652532	-2.525994	-0.115608	-0.552955
38	-1.177178	-1.997270	-0.212785	-0.771497
39	-0.773820	-1.383122	-0.297583	-0.777094
40	-0.989577	-0.100900	-0.315786	-0.846411
41	-1.008534	-1.320796	-0.384552	-0.857466
42	-1.102106	-0.906938	-0.520596	-1.015365
43	-1.281134	0.217682	-1.449605	-1.058960
44	-1.134305	1.206419	-1.509074	-1.173209
45	-1.600350	0.101254	-1.727400	-1.180082
46	-1.593413	-0.199322	0.711122	-1.566922
47	-1.622362	0.507722	-1.743127	-1.740627
48	-1.610433	-2.509409	-1.743127	-1.913212
49	-1.622362	-0.157226	-1.369985	-2.439313

In [36]:

```
model6 = smf.ols("Profit~RDS+Administration+Marketing",data=data).fit()  
model6.summary()
```

Out[36]:

OLS Regression Results			
Dep. Variable:	Profit	R-squared:	0.951
Model:	OLS	Adj. R-squared:	0.948
Method:	Least Squares	F-statistic:	296.0
Date:	Sun, 10 Apr 2022	Prob (F-statistic):	4.53e-30
Time:	19:47:32	Log-Likelihood:	4.3222
No. Observations:	50	AIC:	-0.6444
Df Residuals:	46	BIC:	7.004
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-4.927e-16	0.033	-1.51e-14	1.000	-0.066	0.066
RDS	0.9176	0.051	17.846	0.000	0.814	1.021

Marketing 0.0826 0.050 1.655 0.105 -0.018 0.183

Omnibus:	14.838	Durbin-Watson:	1.282
Prob(Omnibus):	0.001	Jarque-Bera (JB):	21.442
Skew:	-0.949	Prob(JB):	2.21e-05
Kurtosis:	5.586	Cond. No.	2.78

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [37]: model7 = smf.ols("Profit~RDS",data=data).fit()  
model7.summary()
```

Out[37]:

OLS Regression Results			
Dep. Variable:	Profit	R-squared:	0.947
Model:	OLS	Adj. R-squared:	0.945
Method:	Least Squares	F-statistic:	849.8
Date:	Sun, 10 Apr 2022	Prob (F-statistic):	3.50e-32
Time:	19:47:33	Log-Likelihood:	2.2714
No. Observations:	50	AIC:	-0.5428
Df Residuals:	48	BIC:	3.281
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-5.274e-16	0.033	-1.58e-14	1.000	-0.067	0.067
RDS	0.9729	0.033	29.151	0.000	0.906	1.040

Omnibus:	13.727	Durbin-Watson:	1.116
Prob(Omnibus):	0.001	Jarque-Bera (JB):	18.536
Skew:	-0.911	Prob(JB):	9.44e-05
Kurtosis:	5.361	Cond. No.	1.00

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [38]: model8=smf.ols('Profit~Administration',data=data).fit()  
model8.summary()
```

Out[38]:

OLS Regression Results			
Dep. Variable:	Profit	R-squared:	0.040
Model:	OLS	Adj. R-squared:	0.020
	Least Squares	F-statistic:	2.015

Date:	Sun, 10 Apr 2022	Prob (F-statistic):	0.162
Time:	19:47:33	Log-Likelihood:	-69.919
No. Observations:	50	AIC:	143.8
Df Residuals:	48	BIC:	147.7
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-5.274e-16	0.141	-3.73e-15	1.000	-0.284	0.284
Administration	0.2007	0.141	1.419	0.162	-0.084	0.485

Omnibus:	0.126	Durbin-Watson:	0.099
Prob(Omnibus):	0.939	Jarque-Bera (JB):	0.110
Skew:	0.093	Prob(JB):	0.947
Kurtosis:	2.866	Cond. No.	1.00

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [39]:

```
model8=smf.ols("Profit~Marketing",data=data).fit()
model8.summary()
```

Out[39]:

OLS Regression Results

Dep. Variable:	Profit	R-squared:	0.559
Model:	OLS	Adj. R-squared:	0.550
Method:	Least Squares	F-statistic:	60.88
Date:	Sun, 10 Apr 2022	Prob (F-statistic):	4.38e-10
Time:	19:47:33	Log-Likelihood:	-50.470
No. Observations:	50	AIC:	104.9
Df Residuals:	48	BIC:	108.8
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-5.274e-16	0.096	-5.5e-15	1.000	-0.193	0.193
Marketing	0.7478	0.096	7.803	0.000	0.555	0.940

Omnibus:	4.420	Durbin-Watson:	1.178
Prob(Omnibus):	0.110	Jarque-Bera (JB):	3.882
Skew:	-0.336	Prob(JB):	0.144
Kurtosis:	4.188	Cond. No.	1.00

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Calculating VIF

In [40]:

```
rsq_RandD = smf.ols('RDS~Administration+Marketing', data=data).fit().rsquared
vif_RandD = 1/(1-rsq_RandD)

rsq_admin = smf.ols('Administration~RDS+Marketing', data=data).fit().rsquared
vif_admin = 1/(1-rsq_admin)

rsq_marketing = smf.ols('Marketing~RDS+Administration', data=data).fit().rsquared
vif_marketing = 1/(1-rsq_marketing)

# Storing vif values in a data frame
d1 = {'Variables': ['RDS', 'Administration', 'Marketing'], 'VIF': [vif_RandD, vif_admin, vif_marketing]}
Vif_frame = pd.DataFrame(d1)
Vif_frame
```

Out[40]:

	Variables	VIF
0	RDS	2.468903
1	Administration	1.175091
2	Marketing	2.326773

Observations

The vif is less hence as a result there is no colinearity in the features


Residual Analysis

In [41]:

```
import statsmodels.api as sm
```

In [42]:

```
qqplot = sm.qqplot(model5.resid, line="q")
plt.title("Normal Q-Q Plot of Residuals")
plt.show()
```



Observation:


Errors are coming from normal distribution

Residuals For Homoscedasticity

```
In [43]: def get_standardized_values( vals ):
          return (vals - vals.mean())/vals.std()
```

```
In [44]: plt.scatter(get_standardized_values(model5.fittedvalues),
                    get_standardized_values(model5.resid))

plt.title('Residual Plot')
plt.xlabel('Standardized Fitted values')
plt.ylabel('Standardized residual values')
plt.show()
```

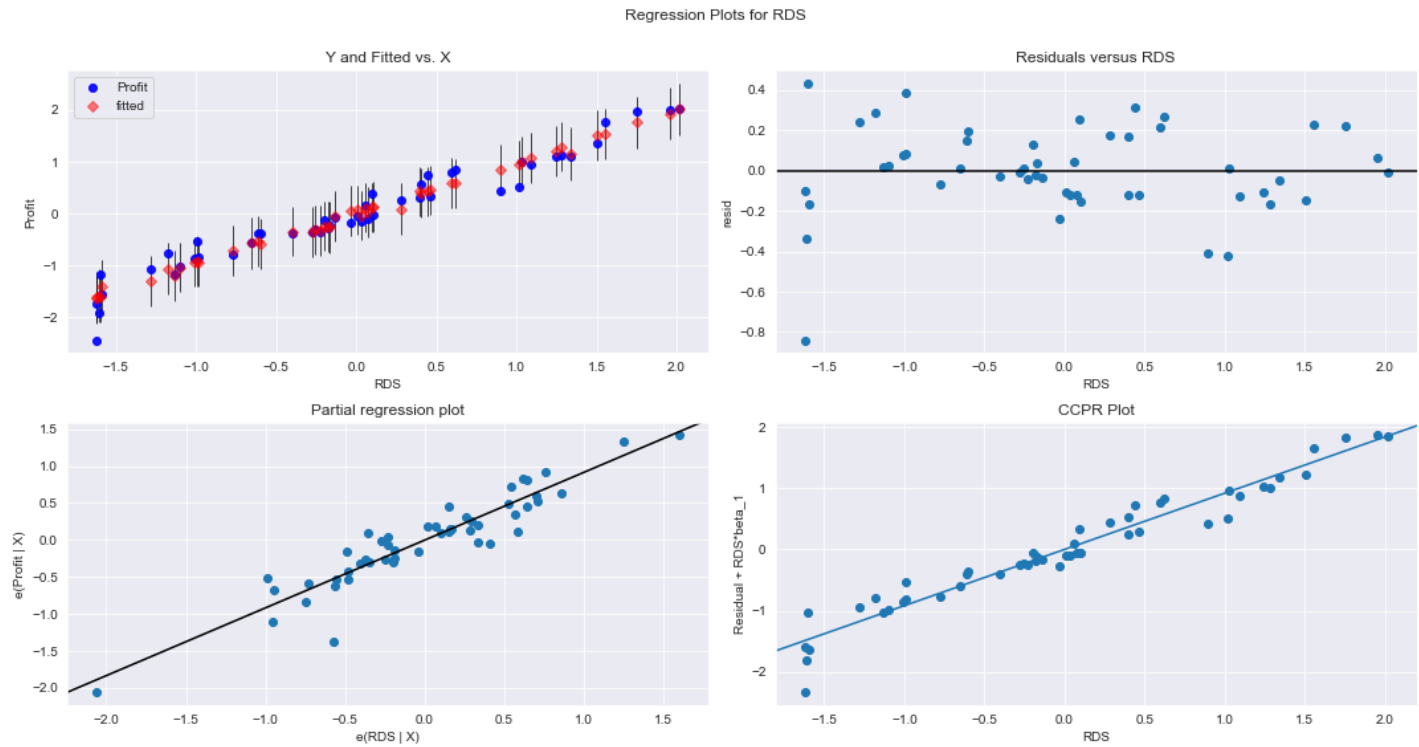


Observation:

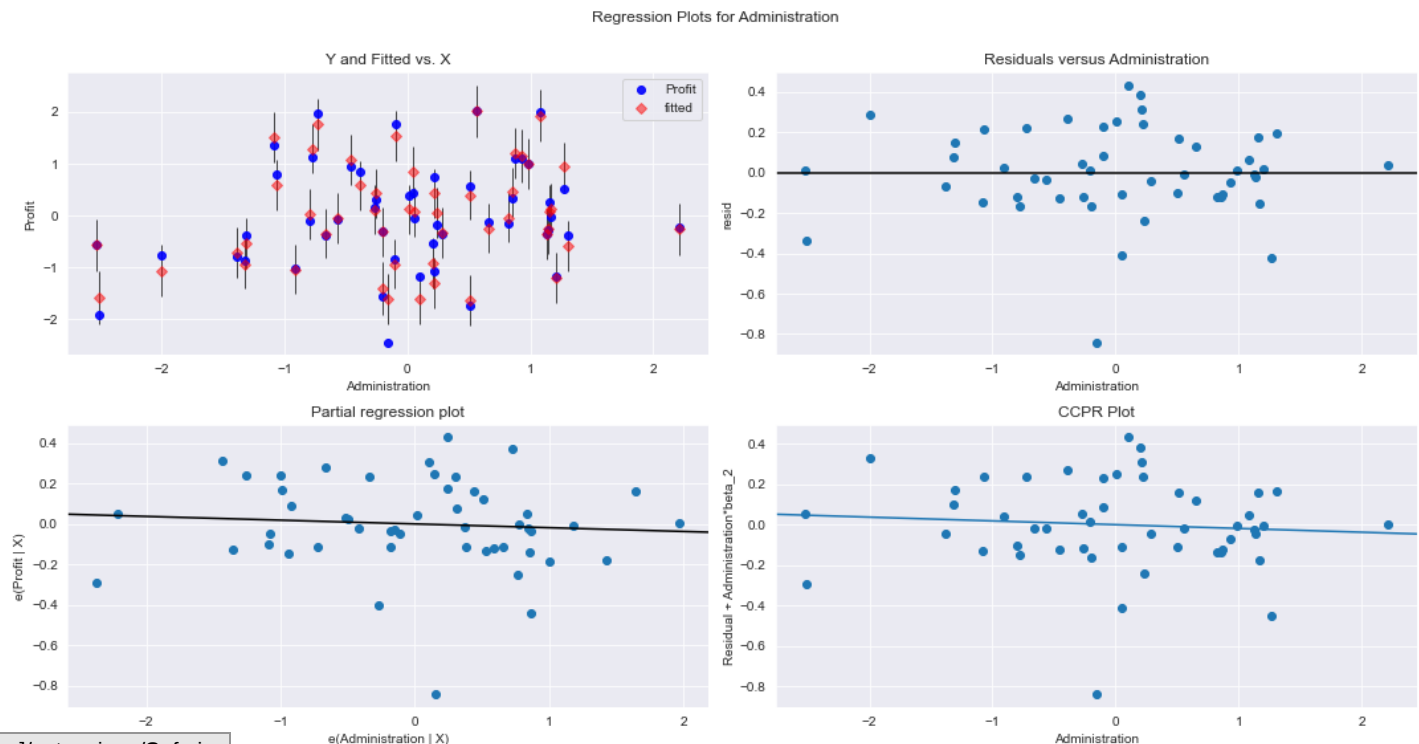
So we can see that there is no such pattern is creating and the data is randomly scattered so this data is an homoscedasticity

Residual vs Regressor

```
In [45]: fig = plt.figure(figsize=(15,8))
fig = sm.graphics.plot_regress_exog(model6, "RDS", fig=fig)
plt.show()
```

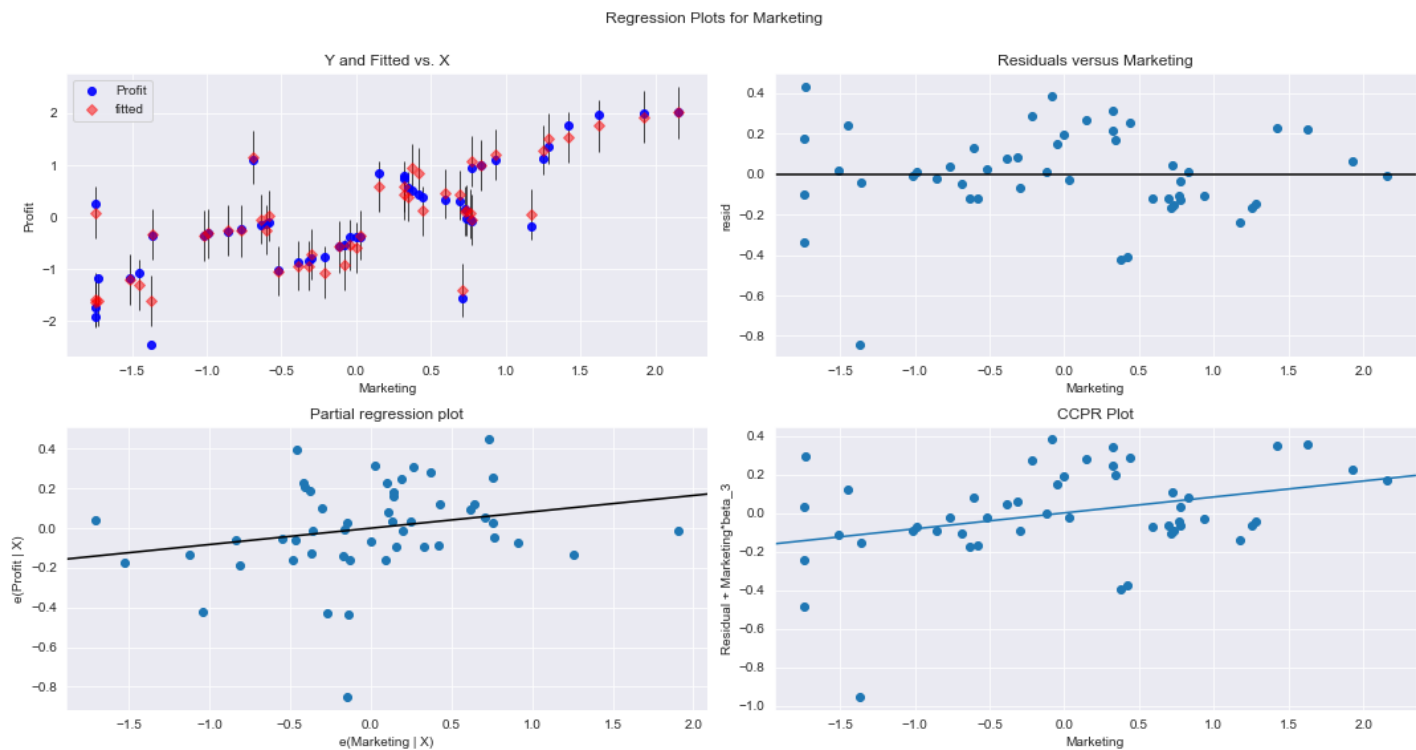


```
In [46]: fig = plt.figure(figsize=(15,8))
fig = sm.graphics.plot_regress_exog(model6, "Administration", fig=fig)
plt.show()
```



In [47]:

```
fig = plt.figure(figsize=(15,8))
fig = sm.graphics.plot_regress_exog(model6, "Marketing",fig=fig)
plt.show()
```




Outlier Detection

In [48]:

```
model5_influence=model5.get_influence()
(c,_)=model5_influence.cooks_distance
```


In [49]:

```
fig = plt.subplots(figsize=(20, 7))
plt.stem(np.arange(len(df1)), np.round(c, 3))
plt.xlabel('Row index')
plt.ylabel('Cooks Distance')
plt.show()
```



In [50]:

```
from statsmodels.graphics.regressionplots import influence_plot
influence_plot(model5)
```



```
In [51]: k = df1.shape[1]
n = df1.shape[0]
leverage_cutoff = 3*((k + 1)/n)
leverage_cutoff
```

```
Out[51]: 0.30000000000000004
```

```
In [52]: (np.argmax(c), np.max(c))
```

```
Out[52]: (19, 0.18507855145120508)
```

```
In [53]: data[df.index.isin([19, 45, 48, 49])]
```

```
Out[53]:
```

	RDS	Administration	Marketing	Profit
19	0.279442	1.159837	-1.743127	0.269773
45	-1.600350	0.101254	-1.727400	-1.180082
48	-1.610433	-2.509409	-1.743127	-1.913212
49	-1.622362	-0.157226	-1.369985	-2.439313

```
In [54]: data2=data.copy()
data2
```

```
Out[54]:
```

	RDS	Administration	Marketing	Profit
0	2.016411	0.560753	2.153943	2.011203
1	1.955860	1.082807	1.923600	1.999430
2	1.754364	-0.728257	1.626528	1.980842
3	1.554784	-0.096365	1.422210	1.776627
4	1.504937	-1.079919	1.281528	1.357740
5	1.279800	-0.776239	1.254210	1.127250
6	1.340066	0.932147	-0.688150	1.105481
7	1.279800	0.871980	0.932186	1.096210

	RDS	Administration	Marketing	Profit
8	1.030369	0.986952	0.830887	1.007470
9	1.091819	-0.456640	0.776107	0.946022
10	0.620398	-0.387599	0.149807	0.854847
11	0.593085	-1.065540	0.319834	0.808168
12	0.443260	0.215449	0.320617	0.741155
13	0.402078	0.510179	0.343957	0.558750
14	1.017181	1.269199	0.375742	0.516026
15	0.897913	0.045868	0.419219	0.448720
16	0.094441	0.009118	0.440446	0.375436
17	0.460720	0.855666	0.591017	0.334771
18	0.396725	-0.258465	0.692992	0.307116
19	0.279442	1.159837	-1.743127	0.269773
20	0.055726	-0.269588	0.723926	0.161935
21	0.102724	1.169186	0.732788	-0.017534
22	0.006007	0.051850	0.762376	-0.041613
23	-0.136201	-0.562211	0.774349	-0.082169
24	0.073115	-0.795469	-0.581939	-0.086729
25	-0.199312	0.656489	-0.603517	-0.115493
26	0.035370	0.821718	-0.635835	-0.157367
27	-0.035519	0.235069	1.174271	-0.175542
28	-0.168793	2.210141	-0.767189	-0.218798
29	-0.178609	1.142457	-0.858134	-0.275882
30	-0.258074	-0.205629	-0.990357	-0.302625
31	-0.276958	1.130554	-1.014419	-0.364127
32	-0.226949	0.283924	-1.362450	-0.365524
33	-0.401129	-0.659324	0.029817	-0.381787
34	-0.600682	1.310535	-0.001879	-0.383444
35	-0.609750	-1.308658	-0.045493	-0.389291
36	-0.991570	0.205925	-0.081763	-0.533932
37	-0.652532	-2.525994	-0.115608	-0.552955
38	-1.177178	-1.997270	-0.212785	-0.771497
39	-0.773820	-1.383122	-0.297583	-0.777094
40	-0.989577	-0.100900	-0.315786	-0.846411
41	-1.008534	-1.320796	-0.384552	-0.857466
42	-1.102106	-0.906938	-0.520596	-1.015365
43	-1.281134	0.217682	-1.449605	-1.058960
44	-1.134305	1.206419	-1.509074	-1.173209
45	-1.600350	0.101254	-1.727400	-1.180082
46	-1.593413	-0.199322	0.711122	-1.566922

	RDS	Administration	Marketing	Profit
47	-1.622362	0.507722	-1.743127	-1.740627
48	-1.610433	-2.509409	-1.743127	-1.913212
49	-1.622362	-0.157226	-1.369985	-2.439313

In [56]:

```
data2=data.drop(data.index[[19,45,48,49]],axis=0).reset_index(drop=True)
data2
```

Out[56]:

	RDS	Administration	Marketing	Profit
0	2.016411	0.560753	2.153943	2.011203
1	1.955860	1.082807	1.923600	1.999430
2	1.754364	-0.728257	1.626528	1.980842
3	1.554784	-0.096365	1.422210	1.776627
4	1.504937	-1.079919	1.281528	1.357740
5	1.279800	-0.776239	1.254210	1.127250
6	1.340066	0.932147	-0.688150	1.105481
7	1.245057	0.871980	0.932186	1.096210
8	1.030369	0.986952	0.830887	1.007470
9	1.091819	-0.456640	0.776107	0.946022
10	0.620398	-0.387599	0.149807	0.854847
11	0.593085	-1.065540	0.319834	0.808168
12	0.443260	0.215449	0.320617	0.741155
13	0.402078	0.510179	0.343957	0.558750
14	1.017181	1.269199	0.375742	0.516026
15	0.897913	0.045868	0.419219	0.448720
16	0.094441	0.009118	0.440446	0.375436
17	0.460720	0.855666	0.591017	0.334771
18	0.396725	-0.258465	0.692992	0.307116
19	0.055726	-0.269588	0.723926	0.161935
20	0.102724	1.169186	0.732788	-0.017534
21	0.006007	0.051850	0.762376	-0.041613
22	-0.136201	-0.562211	0.774349	-0.082169
23	0.073115	-0.795469	-0.581939	-0.086729
24	-0.199312	0.656489	-0.603517	-0.115493
25	0.035370	0.821718	-0.635835	-0.157367
26	-0.035519	0.235069	1.174271	-0.175542
27	-0.168793	2.210141	-0.767189	-0.218798
28	-0.178609	1.142457	-0.858134	-0.275882
29	-0.258074	-0.205629	-0.990357	-0.302625
30	-0.276958	1.130554	-1.014419	-0.364127
		0.283924	-1.362450	-0.365524

	RDS	Administration	Marketing	Profit
32	-0.401129	-0.659324	0.029817	-0.381787
33	-0.600682	1.310535	-0.001879	-0.383444
34	-0.609750	-1.308658	-0.045493	-0.389291
35	-0.991570	0.205925	-0.081763	-0.533932
36	-0.652532	-2.525994	-0.115608	-0.552955
37	-1.177178	-1.997270	-0.212785	-0.771497
38	-0.773820	-1.383122	-0.297583	-0.777094
39	-0.989577	-0.100900	-0.315786	-0.846411
40	-1.008534	-1.320796	-0.384552	-0.857466
41	-1.102106	-0.906938	-0.520596	-1.015365
42	-1.281134	0.217682	-1.449605	-1.058960
43	-1.134305	1.206419	-1.509074	-1.173209
44	-1.593413	-0.199322	0.711122	-1.566922
45	-1.622362	0.507722	-1.743127	-1.740627

In [57]: `data2.shape`

Out[57]: (46, 4)

Improving Model

In [58]: `final_model= smf.ols('Profit~RDS+Administration+Marketing',data=data2).fit()`

In [59]: `final_model.summary()`

Out[59]:

OLS Regression Results											
Dep. Variable:	Profit		R-squared:	0.965							
Model:	OLS		Adj. R-squared:	0.962							
Method:	Least Squares		F-statistic:	384.0							
Date:	Sun, 10 Apr 2022		Prob (F-statistic):	1.53e-30							
Time:	19:51:40		Log-Likelihood:	15.919							
No. Observations:	46		AIC:	-23.84							
Df Residuals:	42		BIC:	-16.52							
Df Model:	3										
Covariance Type:	nonrobust										
	coef	std err	t	P> t 	[0.025	0.975]					
Intercept	0.0162	0.027	0.606	0.548	-0.038	0.070					
RDS	0.8944	0.042	21.210	0.000	0.809	0.980					
Administration	-0.0447	0.030	-1.504	0.140	-0.105	0.015					
	0.0769	0.044	1.763	0.085	-0.011	0.165					

Omnibus:	0.048	Durbin-Watson:	1.731
Prob(Omnibus):	0.976	Jarque-Bera (JB):	0.243
Skew:	-0.002	Prob(JB):	0.886
Kurtosis:	2.644	Cond. No.	2.76

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

MSE

```
In [60]: final_model.mse_resid
```

```
Out[60]: 0.032095711114433942
```

RMSE

```
In [61]: np.sqrt(final_model.mse_resid)
```

```
Out[61]: 0.17915275924288585
```

```
In [62]: data3 = sc.inverse_transform(data2)
```

Transforming Standard Scaler Data Into Raw Data

```
In [63]: data4=pd.DataFrame(data3,columns=["RDS", "Administration", "Marketing", "Profit"])
data4
```

```
Out[63]:
```

	RDS	Administration	Marketing	Profit
0	165349.20	136897.80	471784.10	192261.83
1	162597.70	151377.59	443898.53	191792.06
2	153441.51	101145.55	407934.54	191050.39
3	144372.41	118671.85	383199.62	182901.99
4	142107.34	91391.77	366168.42	166187.94
5	131876.90	99814.71	362861.36	156991.12
6	134615.46	147198.87	127716.82	156122.51
7	130298.13	145530.06	323876.68	155752.60
8	120542.52	148718.95	311613.29	152211.77
9	123334.88	108679.17	304981.62	149759.96
10	101913.08	110594.11	229160.95	146121.95
11	100671.96	91790.61	249744.55	144259.40

	RDS	Administration	Marketing	Profit
12	93863.75	127320.38	249839.44	141585.52
13	91992.39	135495.07	252664.93	134307.35
14	119943.24	156547.42	256512.92	132602.65
15	114523.61	122616.84	261776.23	129917.04
16	78013.11	121597.55	264346.06	126992.93
17	94657.16	145077.58	282574.31	125370.37
18	91749.16	114175.79	294919.57	124266.90
19	76253.86	113867.30	298664.47	118474.03
20	78389.47	153773.43	299737.29	111313.02
21	73994.56	122782.75	303319.26	110352.25
22	67532.53	105751.03	304768.73	108733.99
23	77044.01	99281.34	140574.81	108552.04
24	64664.71	139553.16	137962.62	107404.34
25	75328.87	144135.98	134050.07	105733.54
26	72107.60	127864.55	353183.81	105008.31
27	66051.52	182645.56	118148.20	103282.38
28	65605.48	153032.06	107138.38	101004.64
29	61994.48	115641.28	91131.24	99937.59
30	61136.38	152701.92	88218.23	97483.56
31	63408.86	129219.61	46085.25	97427.84
32	55493.95	103057.49	214634.81	96778.92
33	46426.07	157693.92	210797.67	96712.80
34	46014.02	85047.44	205517.64	96479.51
35	28663.76	127056.21	201126.82	90708.19
36	44069.95	51283.14	197029.42	89949.14
37	20229.59	65947.93	185265.10	81229.06
38	38558.51	82982.09	174999.30	81005.76
39	28754.33	118546.05	172795.67	78239.91
40	27892.92	84710.77	164470.71	77798.83
41	23640.93	96189.63	148001.11	71498.49
42	15505.73	127382.30	35534.17	69758.98
43	22177.74	154806.14	28334.72	65200.33
44	1315.46	115816.21	297114.46	49490.75
45	0.00	135426.92	0.00	42559.73

In [64]:

```
final_model2 = smf.ols("Profit~RDS+Administration+Marketing",data=data4).fit()  
final_model2.summary()
```

Out[64]:

OLS Regression Results

Dep. Variable:	Profit	R-squared:	0.965
----------------	--------	------------	-------

Model:	OLS	Adj. R-squared:	0.962
Method:	Least Squares	F-statistic:	384.0
Date:	Sun, 10 Apr 2022	Prob (F-statistic):	1.53e-30
Time:	19:51:53	Log-Likelihood:	-471.41
No. Observations:	46	AIC:	950.8
Df Residuals:	42	BIC:	958.1
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	5.721e+04	5824.368	9.822	0.000	4.55e+04	6.9e+04
RDS	0.7854	0.037	21.210	0.000	0.711	0.860
Administration	-0.0642	0.043	-1.504	0.140	-0.150	0.022
Marketing	0.0253	0.014	1.763	0.085	-0.004	0.054

Omnibus:	0.048	Durbin-Watson:	1.731
Prob(Omnibus):	0.976	Jarque-Bera (JB):	0.243
Skew:	-0.002	Prob(JB):	0.886
Kurtosis:	2.644	Cond. No.	1.60e+06

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.6e+06. This might indicate that there are strong multicollinearity or other numerical problems.

MSE

```
In [65]: final_model2.mse_resid
```

```
Out[65]: 51099466.48746972
```

RMSE

```
In [66]: np.sqrt(final_model2.mse_resid)
```

```
Out[66]: 7148.38908338583
```

Observation :

As We can see that the R-square of the final model has been increased by 96%

Predicting New Values:Manually

```
In [67]: values = pd.DataFrame({"RDS":86000,"Administration":1234567.87,"Marketing":345678.56}, index=values)
```

Out[67]:

	RDS	Administration	Marketing
1	86000	1234567.87	345678.56

```
In [68]: pd.DataFrame(final_model.predict(values), columns=["Profit"])
```

Out[68]:

	Profit
1	48371.332351

Automatic Predictions

```
In [69]: pred_y = final_model2.predict(data4)
```

```
In [70]: prediction_values=pd.DataFrame(pred_y, columns=["Profit"])
prediction_values
```

Out[70]:

	Profit
0	190233.337001
1	186435.222637
2	181559.701604
3	172683.994403
4	172225.976821
5	163566.198942
6	156712.159018
7	158401.007275
8	150223.390224
9	154820.766302
10	135951.487205
11	136706.556594
12	129079.256744
13	127155.952232
14	147853.055789
15	145909.909162
16	117365.840984
17	129391.336010
18	129405.707912
19	117350.723149
20	116491.364017

	Profit
22	111177.321664
23	114901.028567
24	102525.013706
25	110506.847419
26	114577.043666
27	100343.390854
28	101616.554582
29	100777.001811
30	97648.213213
31	99873.617383
32	99610.712750
33	88881.503340
34	93091.317053
35	76654.520276
36	93518.547686
37	73554.385096
38	86594.985022
39	76554.242806
40	77840.473922
41	73346.079715
42	62101.938257
43	65397.637643
44	58330.939210
45	48506.433703

Table of R Square

```
In [71]: R_square={'Prepared_models':['Model', 'Final_Model'], 'R_Squared':[model1.rsquared, final_model1.rsquared]}
table=pd.DataFrame(R_square)
table
```

```
Out[71]:
```

	Prepared_models	R_Squared
0	Model	0.950746
1	Final_Model	0.964824

```
In [ ]:
```